Algorithms

Eyal Arubas

February 11, 2013

> **Warning:**
> This is a rough, incomplete and inaccurate (probably with many typos) draft.
> Use at your own risk.

This notebook is based on an algorithms course I took in 2012 at the Hebrew University of Jerusalem, Israel. The material is based on my notes from the lectures of Prof. Alex Samorodnitsky, as well as some entries in Wikipedia and more.

I wrote this notebook because I find the subject interesting, and it helped me prepare for my exam. Hopefully it will help whoever is reading it as well.

Needless to say, I take no responsibility for the accuracy, completeness and correctness of what is written here. I'm not, in any way, an authority on algorithms, so take it as it is. That being said, I still wish this book to be as helpful as possible, so if you find any mistakes or inaccuracies, please send me an email to **EyalArubas@gmail.com**.

Also, you are more than welcomed to just tell me what you think of this book. I like the feedback.

The structure of the book is such that each chapter will begin with several examples of relevant problems. These examples will demonstrate a situation in which we need to reach a solution by solving a problem of a certain type. Then in the rest of the chapter we will discuss methods for solving this type of problems, as well as many examples. Each chapter is independent, so you can just jump directly to a subject of your choice.

The latest version of this book can be downloaded from my website at **http://EyalArubas.com/AlgorithmsBook**.

I encourage you to share this book and pass it along, if you find it useful of course.

Contents

Transversality is the characterization of intersection between sets.

Chapter 1

Greedy Algorithms

1.1 Exchange Lemmas

Oftentimes we can formulate an algorithm which we theorize solves some optimization problem. If it indeed does solve the problem, we want to be able to prove it. Usually we do it by saying that if some optimal solution to the problem exists, then it will coincide with the solution of our algorithm. Alternatively, we can say that given any other solution to the problem, our solution will be better. Our goal, then, is to show that we can take the other (optimal or not) solution (which is usually not unique), manipulate it without damaging its optimality, if it's optimal, or make it better if it's not optimal; and reach the solution given by our algorithm; thus showing that it, too, is an optimal solution.

For this purpose we use *exchange lemmas*. What these lemmas actually do is show that we can manipulate the uknown solution, i.e. *exchange* part of it with part of our own solution. If we can show that eventually we can replace the **entire** uknown solution with our own solution, and be optimal, then we have shown that our solution is as good and as optimal as any other.

Unknown optimal solution sounds like a fallacy. How can a solution be optimal without knowing what it is? The fact is that we don't need to know exactly what this optimal solution is, but we just need to know what characteristics it must hold in order to be optimal.

To demonstrate, we show several examples.

1.1.1 Fractional Knapsack Problem

This is a similar problem to the "Robbing a bank" example given at **??**. We have a knapsack which can carry a certain amount of weight. We also have a list of items; each item has it's own weight and value. We want to insert items into the knapsack such that the total value of items inside the knapsack is maximal. Notice that items needn't be whole, and can be inserted partially into the knapsack.

Our input is:

1. W - The maximum wight the knapsack can carry.

2. A list of n items. Item i is represented by the pair (v_i, w_i), where v_i is the value of the item and w_i is the weight of the item. All values and wights are non negative.

Our output should be:

A list of numbers $x_1, x_2, ..., x_n$, where x_i is the fractional amount of item i which is inserted into the knapsack ($0 \leq x_i \leq 1$).

The numbers x_i adhere to the constraint $\sum_{i=1}^{n} x_i w_i \leq W$.

Our goal is to maximize the total value of items in the knapsack $\sum_{i=1}^{n} x_i v_i$.

1.1.1.1 Algorithm

We propose a greedy algorithm which yields an optimal solution to this problem. We notice that greediness is the most natural approach in this case, since our goal is to maximal the value of the knapsack, and we can (intuitivly) achieve that by grabing as much as possible from the most valued items.

[TODO]

1.1.2 Independent Vectors Set Problem

Suppose we have a finite vectors set F of n vectors, in some vector space V, and a positive weight functions on these vectors $\mu : V \to R^+$ (this function assigns a positive scalar value to each vector).

Our goal is to find a subset S of F ($S \subseteq F$), such that the vectors in S are linearly independent and the total weight of S is maximized ($\mu(S) = \sum_{v \in S} \mu(v)$).

1.1.2.1 Algorithm

In this problem, too, it's evident that greediness is the most intuitive approach. We want to maximize the weight of S, so we just add the vectors with the maximal weight, as long as linear independence in S is preserved.

Formally, our algorithm is:

Algorithm 1.1 Maximal weight independent vectors set algorithm

---

1. Sort the vectors in F by thier weight in descending order, such that $\mu(v_1) \geq \mu(v_2) \geq ... \geq \mu(v_n)$

2. Initialize S as the empty set: $S = \emptyset$

3. For $i = 1...n$:

    (a) If $S \cup \{v_i\}$ is linearly independent, update: $S = S \cup \{v_i\}$

4. Return S

1.1.2.2 Proof

We want to prove that algorithm 1.1 indeed returns a set of linearly independent vectors with maximal weight.

For this, we need to use the following lemma:

Lemma 1. *Let A, B two finite subsets of linearly independent vectors in vector space V. Suppose $|A| > |B|$. Then there is a vector $v \in A \backslash B$ such that $B \cup \{v\}$ in linearly independent.*

In others words, if A has more vectors than B, then we can find some vector v in A and add it to B such that $B \cup \{v\}$ is also linearly independent.

Proof. We need to show that there is a $v \in A \backslash B$ that is linearly independent with B. We will prove by negation. Let's suppose that there is no such vector. This means that **all** vectors in A are linearly dependent with B. In other words, $\forall v \in A : v \in B$, which means that $span\,(A) \subseteq span\,(B)$. But then $\dim\,(span\,(A)) \leq \dim\,(span\,(B))$.

Because A and B are sets of linearly independent vectors, then $|A| = \dim\,(span\,(A))$ and $|B| = \dim\,(span\,(B))$, but this means that $|A| \leq |B|$, which is a contradiction to our assumption that $|A| > |B|$. Thus we conclude that indeed there is a vector $v \in A \backslash B$ such that $B \cup \{v\}$ in linearly independent. $\square$

Now we can continue with the proof of the optimality of algorithm 1.1.

Remember that we want to prove that the set S which is returned by our algorithm has the maximal weight of all linearly independent subsets of F.

We will prove by negation.

Suppose there is some better, optimal, solution T such that $\mu\,(T) > \mu\,(S)$.

By lemma 1, $|S| = |T|$, because:

1. If $|S| > |T|$, then by the lemma, there is a vector $v \in S \backslash T$ such that $T \cup \{v\}$ is linearly independent. But this contradict the optimality of T, thus it's impossible.

2. If $|T| > |S|$, then by the lemma, there is a vector $v \in T \backslash S$ such that $S \cup \{v\}$ is linearly independent. But this contradicts the operation of our algorithm, which was supposed to add this vector v to S. So this is also impossible.

Thus indeed $|S| = |T|$.

Let's write the vectors in S and T by descending weight order:

$S = \{v_1, v_2, ..., v_k\} \quad \mu\,(v_1) \geq \mu\,(v_2) \geq ... \geq \mu\,(v_k)$
$T = \{u_1, u_2, ..., u_k\} \quad \mu\,(u_1) \geq \mu\,(u_2) \geq ... \geq \mu\,(u_k)$

Because $\mu\,(T) > \mu\,(S)$, then there must be some index i which is the first occurance of $\mu\,(u_i) > \mu\,(v_i)$.

We denote:

$A = \{v_1, ..., v_{i-1}\}$
$B = \{u_1, ..., u_{i-1}, u_i\}$

Obviously both A and B are sets of linearly independent vectors, and $|B| >$ $|A|$. By lemma 1, there is a vector $u \in B \backslash A$ such that $A \cup \{u\}$ is linearly independent. Because S and T are ordered by descending weights, then the weight of this vector u (whichever it may be) is at least as the smallest-weight vector in B, which is u_i. In other words $\mu(u) \geq \mu(u_i)$. And by the definition of i, also $\mu(u_i) > \mu(v_i)$, and thus $\mu(u) > \mu(v_i)$.

Finally we notice that this contradicts the operation of our algorithm. Because, if $\mu(u) > \mu(v_i)$, then our algorithm should've chosen u before v_i (as we have seen, $A \cup \{u\}$ is linearly independent).

We have reached a contradiction, which finishes our proof. ∎

1.2 Matroids

1.2.1 Properties of a Matroid

We define matroids through their properties. A matroid M is a pair of two entities:

1. A set, S, of elements.

2. A set, I, of subsets of S. More specifically, I is a subset of the power set* of S: $I \subseteq P(S)$.

We denote the matroid which is composed of S and I with $M = (S, I)$. We require I to hold three properties:

1. Must contain the empty set ($\emptyset \in I$).

2. Must hold the *Hereditary Property*.

3. Must hold the *Augmentation Property*.

1.2.1.1 Heriditary Property

We remember that I is a set of subsets of S. In other words, the elements of I are sets themselfs. If we take some set A in I ($A \in I$), then we can also look at subsets of A. Let's look at some subset of A and denote it with B, so we have $B \subseteq A$. Now we can ask an interesting question - does $B \in I$ hold? In other words, is B also an element in I? The answer is that sometimes it is and sometimes it isn't. But if it is, for all such A's and B's, then we say that I is *heriditary*.

Formal Definition I is heriditary if and only if $\forall A \in I, \forall B \subseteq A : B \in I$.
And in words - For every set A in I, every subset of A is also in I.

---

*The power set, $P(S)$, of a set S, is defined to be the set of all the subsets of S. The power set is usually denoted by $P(S)$ or by 2^S. For example, given a set $S = \{1, 2\}$, then $P(S) = \{\emptyset, \{1\}, \{2\}, \{1, 2\}\}$.

1.2.1.2   Augmentation Property

Suppose we have two subsets of S which are in I: $A, B \in I$. And suppose that B contains more elements than A. Then is it possible to take some element from B, which is not also in A, move it to A (**augment** A), and still remain in I? If it's possible for all A's and B's in I, then we say that I has the augmentation property.

Formal Definition I has the augmentation property if and only if $\forall A, B \in I, |B| > |A|, \exists x \in B \backslash A : A \cup \{x\} \in I$.

And in words - For all sets A and B in I, such that B is bigger than A, there exists an element x which is in B but not in A, such that the union of A and x is also in I.

Note: We only have to find just one such x.

1.2.2   The Generic Greedy Algorithm

A generic greedy algorithm will solve an optimization problem if it is defined as a matroid.

[TODO]

1.2.3   Examples

1.2.3.1   The Transversal Matroid

Let $C_1, C_2, ..., C_k$ be k sets of integers which comprise a disjoint union of $[n]^*$: $C_1 \uplus C_2 \uplus ...C_k = [n]$.

We denote $S = [n]$, and we define $I = \{A \subseteq S : \forall i = 1, ..., k \, |A \cap C_i| \leq c\}$. I Contains all the subsets of S that have an intersection with each C_i of at most c elements.

Transversality is the characterization of intersection between sets.

Claim: $M = (S, I)$ is a matroid.

Proof: We need to prove that I holds each of the required three properties.

Contains the empty set: Suppose $A = \emptyset$. Obviously $A \subset S$ and for each $i = 1, ..., k$ it is evident that $|A \cap C_i| = |\emptyset \cap C_i| = 0 \leq c$. Thus $A \in I$. $\blacksquare$

Hereditary Property: We need to prove that $\forall A \in I, \forall B \subseteq A : B \in I$. Suppose $A \in I$ and $B \subseteq A$.

We need to show that $|B \cap C_i| \leq c$.

We know that $B \subseteq A$, so obviously $|B| \leq |A|$, thus it is evident that $|B \cap C_i| \leq |A \cap C_i|$ for every i. But because $A \in I$, we have $|A \cap C_i| \leq c$, which means $|B \cap C_i| \leq |A \cap C_i| \leq c$ for every i. Thus $B \in I$. $\blacksquare$

$^*[n] = \{1, 2, 3, ..., n\}$

Augmentation Property: We need to prove that $\forall A, B \in I, |B| > |A|, \exists x \in B \backslash A : A \cup \{x\} \in I$.

Suppose $A, B \in I$ and $|B| > |A|$. We need to show that there is $x \in B \backslash A$ such that $A \cup \{x\} \in I$.

[TODO]

Chapter 2

Dynamic Algorithms

Chapter 3

Approximation Algorithms

3.1 Examples

3.1.1 Parallel Machine Online Scheduling

We have tasks we need to perform, and several machines to perform those tasks. We want to assign tasks to machines in such an order that we finish all the tasks as early as possible. By "online" we mean we cannot control the order in which the tasks arrive.

Input:

- k - Number of machines we have.

- $t_1, t_2, ..., t_n$ - n tasks, task i takes t_i amount of time to finish.

Output: A function S which assigns tasks to machines $S : \{t_1, t_2, ..., t_n\} \to \{1, 2, ..., k\}$.

Goal: To minimize $q(S)$ - The time in which the last task finishes.

We will show a greedy algorithm which gives a $\left(2 - \frac{1}{k}\right)$-approximation to the problem.

3.1.1.1 Algorithm

Algorithm 3.1 Parallel machine online scheduling $\left(2 - \frac{1}{k}\right)$-approximation algorithm

1. For each task i from 1 to n:

    (a) Send task i to the machine which is currently planned to finish first.

In other words, the algorithm sends the current task to the least occupied machine.

For example, given the following tasks:
$t_1 = 1, t_2 = \frac{1}{2}, t_3 = \frac{2}{3}, t_4 = 1$
and two machines ($k = 2$), the algorithm will do the following:

1. Task 1 is assigned to machine 1:

Machine 1	Machine 2
t_1	

2. Task 2 is assigned to machine 2:

Machine 1	Machine 2
t_1	t_2

3. Machine 2 is less occupied than machine 1 ($t_2 < t_1$), so task 3 is assigned to machine 2:

Machine 1	Machine 2
t_1	t_2, t_3

4. Machine 1 is less occupied than machine 2 ($t_1 < t_2 + t_3$), so task 4 is assigned to machine 1:

Machine 1	Machine 2
t_1, t_4	t_2, t_3

In this case, machine 1 will need $t_1 + t_4 = 2$ time to finish, and machine 2 will need $t_2 + t_3 = 1\frac{1}{6}$ time to finish. So $q(S) = 2$.

This is an approximating algorithm, so it doesn't promise the best result. And indeed there is a better solution:

Machine 1	Machine 2
t_1, t_2	t_3, t_4

, for which $q(S) = 1\frac{2}{3}$.

3.1.1.2 Proof of $\left(2 - \frac{1}{k}\right)$-approximation

We need to show that the algorithm in 3.1.1.1 is $\left(2 - \frac{1}{k}\right)$-approximating. In other words, if we denote the solution of the algorithm as S, and the uknown optimal solution as $S^\star$, then we need to show that $q(S) \leq \left(2 - \frac{1}{k}\right) q(S^\star)$.

Obviously we don't know what is $S^\star$, but we can determine some characteristics it must hold. We do this in the following two lemmas.

Lemma 2. *The time in which the last machine finishes in the optimal solution, is greater than the average finish time between all machines: $q(S^\star) \geq \frac{1}{k} \sum_{i=1}^{n} t_i$.*

Lemma 3. *The time in which the last machine finishes in the optimal solution, is greater than the time of the longest task (denoted by t_{max}): $q(S^\star) \geq t_{max}$.*

We won't prove these lemmas, as they are self evident. We now continue with the proof.

Let's consider the last decision our algorithm has made regarding the machine which finishes last. We denote the index of the task of this decision as l ($1 \leq l \leq n$), and the machine as j ($1 \leq j \leq k$). We also denote as F_j the time in which machine j finishes after the first $l - 1$ tasks have been assigned.

We also know that, by the definition of the algorithm, $F_j \leq \frac{1}{k} \sum_{i=1}^{l-1} t_i$ (because otherwise task l wouldn't be assigned to this machine).

Since j is the machine which finished last, we know that $q(S) = F_j + t_l$.

Thus we have $q(S) \leq \frac{1}{k} \sum_{i=1}^{l-1} t_i + t_l = \frac{1}{k} \sum_{i=1}^{l} t_i + \left(1 - \frac{1}{k}\right) t_l \leq \frac{1}{k} \sum_{i=1}^{n} t_i + \left(1 - \frac{1}{k}\right) t_{max} \leq q(S^\star) + \left(1 - \frac{1}{k}\right) q(S^\star) = \left(2 - \frac{1}{k}\right) q(S^\star).$ ∎

3.1.2 Set Cover Problem

We have m subsets of $[n]^*$: $A_i \subseteq [n]$ and $1 \leq i \leq m$. We also know that the sets A_i cover all of $[n]$: $\cup_{i=1}^{m} A_i = [n]$.

We want to find the minimal partial group of these sets that still covers all of $[n]$. In other words, we want to find $S \subseteq \{1, ..., m\}$ such that $\cup_{i \in S} A_i = [n]$ and $q(S) = |S|$ is minimal.

3.1.2.1 $\ln(n)$-Approximating algorithm

We propose an iterative algorithm. In each iteration we denote by X the set of numbers in $[n]$ which are still not covered ($X \subseteq [n]$). We also denote by T_i the relevant part of set A_i ($T_i = A_i \cap X$).

Algorithm 3.2 Minimal set cover approximating algorithm

1. Initialize $S = \emptyset$

2. Denote by i the index in which $|T_i|$ is maximal.

3. Update:

   (a) $S = S \cup \{i\}$
   (b) $X = X \backslash T_i$
   (c) $\forall j = 1, ..., m : T_j = T_j \backslash T_i$

4. Repeat steps 2,3 until $X = \emptyset$

5. Return S

3.1.2.2 Proof of $\ln(n)$-approximation

We want to prove that algorithm 3.2 gives a $\ln(n)$-approximation to the optimal solution. In other words, we want to show that $|S| \leq \ln(n) |S^\star|$.

We prove by examining the change in the size of X in each iteration of the algorithm.

We denote n_j the size of X after j iterations. Obvisouly $n_0 = n$ and if $n_j = 0$ then $|S| \leq j$.

$^*[n] = \{1, 2, ..., n\}$

Let's examine the $j+1$ iteration of the algorithm (currently $|S| = j$). As denoted earlier, T_i is the chosen set in this iteration, which means $|T_i|$ is maximal. By our definitions, we have $n_{j+1} = |X \backslash T_i| = n_j - |T_i|$.

We use the following lemma:

Lemma 4. $|T_i| \geq \frac{|X|}{|S^\star|}$.

Proof. We need to show that $|T_i| \geq \frac{|X|}{|S^\star|}$, or equivalently $|X| \leq |S^\star| \cdot |T_i|$.

$|T_i|$ has the maximal size of all T's, so $\forall j \neq i : |T_i| \geq |T_j|$.

$S^\star$ is the optimal solution.

By the definition of X and the T's we can write $|X| = |\cup_{j \in S^\star} T_j|$.

So, $|X| = |\cup_{j \in S^\star} T_j| \leq \cup_{j \in S^\star} |T_j| \leq |S^\star| \cdot |T_i|$, as needed. $\qquad\square$

We continue with the proof of the approximation.

We had $n_{j+1} = n_j - |T_i|$, so by the lemma, $n_{j+1} \leq n_j - \frac{|X|}{|S^\star|} = n_j - \frac{n_j}{|S^\star|} = n_j \left(1 - \frac{1}{|S^\star|}\right)$. This is a recursive formula from which we can conclude:

$n_j \leq n_{j-1} \left(1 - \frac{1}{|S^\star|}\right) \leq n_{j-2} \left(1 - \frac{1}{|S^\star|}\right)^2 \leq \ldots \leq n_0 \left(1 - \frac{1}{|S^\star|}\right)^j = n \left(1 - \frac{1}{|S^\star|}\right)^j$.

We will employ the following lemma (which will not be proven here):

Lemma 5. $(1 + a)^b < e^{ab}$, *for all non zero* $a, b \in \mathbb{R}$.

We use the lemma and obtain: $n_j \leq n \left(1 - \frac{1}{|S^\star|}\right)^j \leq n e^{\frac{-j}{|S^\star|}}$.

Since n_j in an integer, then for $j = |S^\star| \ln(n)$ we obtain $n_j = 0$, which, as stated earlier, means that $|S| \leq |S^\star| \ln(n)$, thus completing the proof. $\blacksquare$

3.1.3 Max-Cut Problem

We have an undirected graph $G = (V, E)$.

A cut in the graph, $C = (A, B)$, is such that $A, B \subseteq V$ and $A \uplus B = V^*$.

We denote with E_C the edges of the cut: $E_C = \{(i, j) \in E : i \in A, j \in B\}$.

Our goal is to find a cut C with maximum edges in the graph (maximal $|E_C|$).

---

*$^*\uplus$ is the symbol for disjoin union, i.e. $A \uplus B = V$ means $A \cup B = V$ and $A \cap B = \emptyset$.

3.1.3.1   2-approximating algorithm

Algorithm 3.3 Max-Cut 2-approximating algorithm

1. Initialize two sets: $A = V$, $B = \emptyset$.

2. For each vertex $i \in V$:

   (a) Denote by X the set to which i belongs (A or B).

   (b) Denote by Y the other set.

   (c) If i has more neighbors in X than in Y, move i to Y.

3. Repeat step 2 until no more changes occur.

At first it may seem the algorithm might never stop. But we notice that the quality of any solution ($|E_C|$) is bounded by the number of edges in the graph ($|E_C| \leq |E|$). In each iteration of step 2, we move a vertex to the set in which it has less neighbors (or leave it, if it's already there), which means that $|E_C|$ either increases or stays the same, but it never decreases. But because $|E_C|$ is bounded, then the algorithm must eventually stop.

3.1.3.2   Proof of 2-approximation

We denote by d_v the degree of vertex v, and by v_C the number of edges in the cut that touch vertex v.

Every edge in the cut touches two vertices, so $|E_C| = \frac{1}{2} \sum_{v \in V} v_C$.

By the operation of the algorithm, it makes sure that for every vertex, at least half of its neighbors are in the other set, which means that at least half of the edges that touch this vetex are in the cut: $\forall v \in V : v_C \geq \frac{1}{2} d_v$.

So: $|E_C| = \frac{1}{2} \sum_{v \in V} v_C \geq \frac{1}{4} \sum_{v \in V} d_v$.

$\sum_{v \in V} d_v$ is exactly twice the edges in the graph (We count the number of edges for every vertex, but every edge touches two vertices): $\sum_{v \in V} d_v = 2 |E|$.

Obviously, the optimal solution (denoted $C^\star$) doesn't have more edges in the cut than there are edges in the graph: $|E_{C^\star}| \leq |E|$

So now we have $|E_C| \geq \frac{1}{4} 2 |E| = \frac{1}{2} |E| \geq \frac{1}{2} |E_{C^\star}|$, which completes the proof. ∎

3.1.4   Max-K-Cut

As in the *Max-Cut* problem in 3.1.3, here also we have an undirected graph $G = (V, E)$. Our goal here is to find a partition of V into k disjoint sets $V_1, V_2, ..., V_k$, such that the sum of the number of edges in the cuts (i.e. between all pairs of V_i, V_j) is maximized.

We notice that the *Max-Cut* problem is a private case of the current problem with $k = 2$.

3.1.4.1 $\left(1 - \frac{1}{k}\right)$-approximating algorithm

This algorithm is very similar to algorithm 3.3. Here, too, in each iteration we move the current vertex into the set in which most of it's neighbors will be from the other sets.

We denote by d_v the degree of vertex v, by N_v^{in} the number of neighbors v has inside it's current set, and by N_v^{out} the number of neighbors v has outside of it's current set.

Algorithm 3.4 Max-K-Cut $\left(1 - \frac{1}{k}\right)$-approximating algorithm

1. Partition V into k disjoint sets $V_1, V_2, ..., V_k$ (randomly).

2. For each vertex $i \in V$:

    (a) If $N_v^{in} > \frac{d_v}{k}$, move v from it's current set V_i to another set V_j, such that the new N_v^{out} is maximal.

3. Repeat step 2 until no more changes occur.

From the same considerations of algorithm 3.3, this algorithm must also eventually stop.

3.1.4.2 Proof of $\left(1 - \frac{1}{k}\right)$-approximation

For a solution S we denote by $|E_S|$ the number of edges in the cuts defined by S. If S is the solution of our algorithm, and $S^\star$ is the optimal solution, we need to show that $|E_S| \geq \left(1 - \frac{1}{k}\right)|E_{S^\star}|$.

We notice the following:

- $|E_S| = \frac{1}{2}\sum_{v \in V} N_v^{out}$

- $\forall v \in V : d_v = N_v^{out} + N_v^{in}$.

- $|E| = \frac{1}{2}\sum_{v \in V} d_v$.

After the algorithm finishes its operation, we know that $\forall v \in V : N_v^{in} < \frac{d_v}{k}$.

Thus we conclude $\forall v \in V : N_v^{out} \geq \frac{k-1}{k}d_v$.

So now we can write: $|E_S| = \frac{1}{2}\sum_{v \in V} N_v^{out} \geq \frac{1}{2}\frac{k-1}{k}\sum_{v \in V} d_v = \frac{k-1}{k}|E|$.

We also know that in any solution, as well as the optimal solution, the number of edges in the cuts cannot exeed the number of edges in the graph, so: $|E_{S^\star}| \leq |E|$.

So finally we have: $|E_S| \geq \frac{k-1}{k}|E| \geq \frac{k-1}{k}|E_{S^\star}| = \left(1 - \frac{1}{k}\right)|E_{S^\star}|$, thus completing the proof. ∎

3.1.5 Vertex Cover

We have a graph of m vertices and n edges: $G = (V, E)$.

We want to find a subset $S \subseteq V$ such that every edge has at least one vertex in that set. In other words $\forall e = (i, j) \in E : i \in S \, or \, j \in S$.

Our goal is to find such a set with minimal size (minimize $q(S) = |S|$).

We notice this is a special case of the *Set Cover Problem* presented in 3.1.2, because if we number the edges of the graph $1, ..., n$, we can define the set A_i to be the vertices of edge i. So we can solve this problem with the same $\ln(n)$-approximating algorithm.

Here, however, we will show a 2-approximating algorithm.

First we will see several examples.

Example 6. The complete graph with m vertices (K_m). In this case $q(S^\star) = m - 1$, because if we leave out 2 vertices, the edge which connects them will not be covered.

Example 7. The bipartite complete graph $K_{a,b}$. If $a \leq b$ then obviously $q(S^\star) = a$.

Example 8. A matching of size k - a graph with k edges where each vertex has degree of excatly 1 (connected to exactly 1 edge). Obviously $q(S^\star) = k$.

Corollary 9. *If a graph **contains** a matching of size k, then $q(S^\star) \geq k$.*

3.1.5.1 2-approximating algorithm

Algorithm 3.5 Vertex cover 2-approximating algorithm

1. Initialize: $S = \emptyset$, $X = E$.

2. While $X \neq \emptyset$:

   (a) Choose some edge $e = (i, j) \in X$.

   (b) Add to S vertices i, j: $S = S \cup \{i, j\}$.

   (c) Remove from X all edges which are covered by i or j: $X = X \setminus \{All\, edges\, that\, touch\, i\, or\, j\}$.

3. Return S.

3.1.5.2 Proof of 2-approximation

We denote the solution of the algorithm with S, and the optimal solution with $S^\star$.

We want to show that $q(S) \leq 2q(S^\star)$ ($|S| \leq 2|S^\star|$).

By the structure of the algorithm, $|S|$ is even (we add two vertices each iteration), so we can write $|S| = 2k$ for some k. We notice that all the edges

the algorithm choses in step 2a have no vertices in common (becuase of step 2c). In other words, these edges make a matching of k edges. By corollary 9 we conclude $|S^\star| \geq k$.

Finally we have $|S| = 2k \leq 2\,|S^\star|$, thus completing the proof. ∎

3.1.6 Metric Travelling Salesman

We have a **complete** undirected graph $G = (V, E)$, and a **metric** weight function $w : E \to \mathbb{R}^+$. Since w is metric, is holds the triangle inequallity: $\forall i, j, k \in V : w\,(i, k) \leq w\,(i, j) + w\,(j, k)$.

Our goal is to find a simple cycle C in the graph $(C = (e_1, e_2, ..., e_k))$ which touches **every** vertex exactly once, and has a minimal weight $(w\,(C) = \sum_{e \in C} w\,(e)$ is minimal).

3.1.6.1 2-approximating algorithm

Algorithm 3.6 Metric Travelling Salesman 2-approximating algorithm

---

1. Initialize $C = \emptyset$.

2. Find a minimal spanning tree (MST) of the graph, and denote it with T.

3. Double the edges in T to form a cycle which touches all the vertices, and denote it with $T^{'}$.

4. Traverse the edges of $T^{'}$ starting from some arbitrary vertex:

   (a) Denote the current vertex with u.

   (b) Denote the next **unvisited** vertex with v.

   (c) Add the edge (u, v) to C: $C = C \cup \{(u, v)\}$.

   (d) Continue the next iteration from v.

5. Return C.

Proof of correctness   We need to prove that the returned C is indeed a simple cycle which touches every vertex exactly once.

We notice that $T^{'}$ is a cycle (not simple) which touches every vertex. In step 4, by traversing **every** edge of $T^{'}$, we visit **every** vertex in the graph. But because we only visit the **unvisited** vertices, we know that we touch every vertex **once**. Because G is a **complete** graph, there is always an edge between the vertices u and v in step 4, thus making C a simple cycle. ∎

3.1.6.2 Proof of 2-approximation

We denote by C the solution of algorithm 3.6, and by $C^\star$ the optimal solution. We need to prove that $w\,(C) \leq 2w\,(C^\star)$.

Claim 10. Let T_{MST} be a minimal spanning tree of G. Then $w(T_{MST}) \leq w(C^\star)$.

Proof. Suppose we have $C^\star$. We remove one of it's edges. Now we have a spanning tree T. Removing an edge from $C^\star$ decreases it's weight by some amount. Thus we have $w(T) \leq w(C^\star)$. From the minimality of T_{MST} we have $w(T_{MST}) \leq w(T)$. So finally we have $w(T_{MST}) \leq w(C^\star)$. ∎ □

In step 4 of the algorithm, from the current vertex u we connect an edge to the next **unvisited** vertex v, so we skip all the visited vertices between them (by the order defined by T'). So by the triangle inequiallity (remember, w is metric), we conclude: $w(C) \leq w\left(T'\right)$.

T' is formed by doubling the edges of T, so $w\left(T'\right) = 2w(T)$.

T is a minimal spanning tree, so by claim 10 we have $w(T) \leq w(C^\star)$.

Finally we can write $w(C) \leq w\left(T'\right) = 2w(T) \leq 2w(C^\star)$. Thus completing the proof. ∎

3.1.6.3   1.5-approximating algorithm (by Christofides)

Algorithm 3.7 Metric Travelling Salesman 1.5-approximating algorithm (by Christofides)

---

1. Initialize $C = \emptyset$.

2. Find a minimal spanning tree for G and denote it with T.

3. Denotations:

   (a) A - The vertices with odd degree in T.
   (b) G_A - The induced graph of A by G.

4. Find a minimal perfect matching M^* in G_A.

5. Let C' be the edges in M^* combined with the edges in T (forming a multigraph).

6. Find an Euler cycle in C' and denote it with C_e.

7. Traverse the edges of C_e starting from some arbitrary vertex:

   (a) Denote the current vertex with u.
   (b) Denote the next **unvisited** vertex with v.
   (c) Add the edge (u, v) to C: $C = C \cup \{(u, v)\}$.
   (d) Continue the next iteration from v.

8. Return C.

Notes:

1. In step 3b, G_A is a complete graph, because it's induced on a complete graph.

2. In step 6, an *Euler cycle* is a cycle (simple or not) which touches every vertex exactly once.

In order to use the algorithm, we need the following claims:

Claim 11. Every graph G has an even number of vertices with odd degree.

Proof. Let d_v denote the degree of vertex v. In every graph $G = (V, E)$, $\sum_{v \in V} d_v = 2|E|$. If the number of vertices with odd d_v is odd, then $\sum_{v \in V} d_v$ will be odd. But this is impossible. ■   □

Claim 12. Every connected graph in which for every vertex v, d_v is even, there is an Euler cycle.

Claim 13. In a graph with even number of vertices, it is possible to find a minimal perfect matching in polynomial runtime (Using Edmond's *Path, Trees and Flowers* algorithm, or Vazirani's algorithm).

Claim 14. In C' there is a cycle.

Proof. We will show that every vertex i in C' has an even degree, so it must contain a cycle. If d_i is even in T, then adding M^* does not affect i's degree, because M^* is a matching of vertices with odd degree. If d_i is odd in T, then adding M^* adds 1 to d_i, because M^* is a matching, and as such connects only one edge to each vertex. $\square$

Claim 15. In C' there is an Euler cycle.

Proof. We have shown that the degree of every vertex in C' is even, so we can use claim 12. $\square$

In step 6 we obtained an Euler cycle, which we transform to a simple cycle which touches every vertex, just as we did in algorithm 3.6.

3.1.6.4 Proof of 1.5-approximation

We denote by C the solution of algorithm 3.7 and by $C^\star$ the optimal solution.

We want to prove $w(C) \le 1.5w(C^\star)$.

Obviously $w(C) \le w(T) + w(M^*)$.

So we will devide the proof into two parts.

First we will show $w(T) \le w(C^\star)$, and then $w(M^*) \le 0.5w(C^\star)$.

From claim 10 we conclude $w(T) \le w(C^\star)$.

Let's denote with $C^\star_A$ the optimal solution for the problem on G_A. By the triangle inequality, we can deduce that an optimal solution which touches only part of the vertices, is better (has less weight) than an optimal solution that touches all of the vertices. Thus we conclude $w(C^\star_A) \le w(C^\star)$.

The set A contains an even number of vertices, and $C^\star_A$ is a cycle which touches all these vertices. This means that in $C^\star_A$ there are two perfect mathces (alternating edges of the cycle). We denote them M^1_A and M^2_A, and assume, WLOG, M^1_A has the smaller weight ($w(M^1_A) \le w(M^2_A)$). Both M^1_A and M^2_A are also perfect matches in G_A, and M^* is defnied to be the **minimal** perfect matching M^* in G_A, thus $w(M^*) \le w(M^1_A)$.

Also, from the definition of M^1_A and M^2_A: $w(C^\star_A) = w(M^1_A) + w(M^2_A)$, and from the minimality of M^1_A: $w(M^1_A) \le \frac{1}{2}w(C^\star_A)$.

From this we conclude $w(M^*) \le w(M^1_A) \le \frac{1}{2}w(C^\star_A) \le \frac{1}{2}w(C^\star)$.

Finally we have $w(C) \le w(T) + w(M^*) \le w(C^\star) + \frac{1}{2}w(C^\star) = 1.5w(C^\star)$, as needed to complete the proof. ∎

3.1.7 3-SAT (Satisfiability)

Before we define the problem, some definitions and denotations:

- Boolean variable: $x \in \{\mathbb{T}, \mathbb{F}\}$.

- Literal: either a variable x or it's negation $\neg x$.

- Clause: A disjunction of literals (such as $x_1 \vee \neg x_2 \vee x_3$).

- Formula: A conjunction of clauses (such as $(x_1 \vee \neg x_2 \vee x_3 \vee x_4) \wedge (\neg x_1 \vee x_4 \vee x_5)$). Such a formula is said to be in *conjunctive normal form* (CNF).

- A *k-CNF* formula is a CNF formula in which each clause contains k literals.

In this problem, our input is a 3-CNF formula with m clauses. Our goal is to find an assignment for the variables in the formula such that the most clauses are satisfied (i.e. evaluated valuated as $\mathbb{T}$).

Let X be an assignment of the variables, we denote $q(X)$ as the number of satisfied clauses.

Assumptions:

1. No clause contains the same literal more than once.

2. No clause contains a variable and it's negation.

3.1.7.1 2-approximating algorithm

Algorithm 3.8 3-SAT 2-approximating algorithm

1. Denote the assignment in which all variables are assigned $\mathbb{T}$ with $X_{\mathbb{T}}$.

2. Denote the assignment in which all variables are assigned $\mathbb{F}$ with $X_{\mathbb{F}}$.

3. If $q(X_{\mathbb{T}}) > q(X_{\mathbb{F}})$, return $X_{\mathbb{T}}$.

4. Else, return $X_{\mathbb{F}}$.

3.1.7.2 Proof of 2-approximation

Denote with m the number of clauses. Let S be the solution of the algorithm and $S^\star$ the optimal solution.

If $X_{\mathbb{T}}$ (as defined in the algorithm) does not satisfy a clause C_i, then it must be satisfied by $X_{\mathbb{F}}$ (because a clause is a disjunction of literals). Thus $q(X_{\mathbb{T}}) + q(X_{\mathbb{F}}) = m$.

We also know that the optimal solution cannot satisfy more than m clauses, so $q(S^\star) \leq m$.

Also, we know that $S = \max\{X_{\mathbb{T}}, X_{\mathbb{F}}\}$, so $q(S) \geq \frac{1}{2}m$.

Finally we have $q(S) \geq \frac{1}{2}m \geq \frac{1}{2}q(S^\star)$, thus completing the proof. ∎

Chapter 4

Probabilistic Algorithms

4.1 Examples

4.1.1 Max-Lin-2

This is a problem of solving a linear system of equations over a finite field of two elements* ($F_2 = \{0, 1\}$).

The input is a system of m equations with n variables over the field F_2.

Our goal is to find an assignment for the n variables, such that as much of the equations are satisfied.

In other words, if we our solution is S and we define $q(S) = number\ of\ satisfied\ equations$, then we would like to maximze $q(S)$.

Notice that if there is an assignment which satisfies **all** the equations, then we already know how to find it by employing one of the methods of solving a linear system of equations. But the more interesting problem in our case, is when not all equations can be satisfied by any assignment. In which case we would like to find the assignment which maximizes the number of satisfied equations.

Example 16. The system of equations $\begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 1 \\ 0 \\ 1 \end{pmatrix}$ is solved with $\begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}$.

Example 17. The system of equations $\begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 1 \\ 0 \\ 1 \\ 1 \end{pmatrix}$ is

---

*$F_2 = \{0, 1\}$ is a finite field (also called *Galois field*), in which the results of arithmetic operations are taken modulus 2. For example: $1 + 1 = 0$.

solved with $\begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}$. Notice that in this case we have more equations than variables.

Example 18. The system of equations $\begin{pmatrix} 1 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 1 \\ 0 \\ 1 \\ 0 \end{pmatrix}$ **can-not** be solved. However, $\begin{pmatrix} x_1 \\ x_2 \\ x_3 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}$ satisfies the top 3 equations, and only the fourth equation is not satisfied.

4.1.1.1 2-approximating probabilistic algorithm

We will now show an algorithm which yields a 2-approximating for the expectancy of the solutions. In other word, if the best solution, $S^\star$, satisfies $q(S^\star)$ equations, then this algorithm finds an assignment, S, such that if we consider the expectancy of these assignments, at least $\frac{1}{2}q(S^\star)$ equations are satisfied ($\mathbb{E}[S] \geq \frac{1}{2}q(S^\star)$).

We denote the n variables with $x_1, x_2, ..., x_n$.

Algorithm 4.1 Max-Lin-2 2-expectancy-approximation probabilistic algorithm

1. For $i = 1, ..., n$:

   (a) Assign $x_i = 0$ or $x_i = 1$ with probability of 0.5 (toss a coin).

4.1.1.2 Proof of 2-approximation of the expectancy

Notice that $q(S^\star) \leq m$, so it's enough to show $\mathbb{E}[S] \geq \frac{m}{2}$.

We define the following probability space:

$\Omega = \{\omega = \{\omega_1, \omega_2, ..., \omega_n\} : \omega_i \in \{0, 1\}\}$

Each $\omega \in \Omega$ is an assignment for $x_1, x_2, ..., x_n$.

We define the probability function $P(\omega) = \frac{1}{2^n}$ (because in each ω we have n elements with probability $\frac{1}{2}$).

We also define a random variable $X(w)$ as the number of equations ω satisfies ($X(\omega) = |\{i : (A\omega)_i = b_i\}|$).

The expectancy of X is the average number of equations satisfied by a random assignment.

We will show that $\mathbb{E}[X] = \frac{m}{2}$.

Let's define m new random variables $X_1(\omega), X_2(\omega), ..., X_m(\omega)$, where $X_i(\omega)$ indicates if equation i is satisfied: $X_i(\omega) = \begin{cases} 1 & (A\omega)_i = b_i \\ 0 & (A\omega)_i \neq b_i \end{cases}$.

By this definition, we have $X(w) = \sum_{i=1}^{m} X_i(\omega)$, and by the linearity of the expectation $\mathbb{E}[X] = \sum_{i=1}^{m} \mathbb{E}[X_i]$.

We will now see that $\forall i = 1, ..., m : \mathbb{E}[X_i] = \frac{1}{2}$, and thus $\mathbb{E}[X] = \frac{m}{2}$.

$$\mathbb{E}[X_i] = \sum_{\omega \in \Omega} P(\omega) \cdot X_i(\omega) = \sum_{\omega \in \Omega : X_i(\omega)=1} P(\omega) = \frac{1}{2^n} \left(\underbrace{\sum_{\omega \in \Omega : X_i(\omega)=1} 1}_{*} \right).$$

The expression marked with $*$ is the number of assignments which can satisfy equation i. From linear algebra we know that a linear equation of n variables has 2^{n-1} legal assignments (over the field F_2)*.

So finally we have $\mathbb{E}[X_i] = \frac{2^{n-1}}{2^n} = \frac{1}{2}$, thus copmleting the proof. ∎

4.1.1.3 Further discussion

Probability of a random assignment We can also determine the probability in which a certain amount of equations are satisfied by an assignment ω.

For example, let's show $P(\omega : X(\omega) \geq 0.4m) \geq \frac{1}{6}$, or in words - we will show that the probability that an assignment ω satisfies at least 40% of the equations is larger than $\frac{1}{6}$.

Let's detone with B the event $B = \{\omega : X(\omega) \geq 0.4m\}$. We want to prove $P(B) \geq \frac{1}{6}$. We will prove by negation.

Suppose $P(B) < \frac{1}{6}$.

$\mathbb{E}[X] = \sum_{\omega \in \Omega} P(\omega) \cdot X(\omega) = \sum_{\omega \in B} P(\omega) \cdot X(\omega) + \sum_{\omega \notin B} P(\omega) \cdot X(\omega)$

Any assignment cannot satisfy more than m equations (obviously, since there are only m equations), so $\sum_{\omega \in B} P(\omega) \cdot X(\omega) \leq \sum_{\omega \in B} P(\omega) \cdot m$.

Also, by the definition of B: $\sum_{\omega \notin B} P(\omega) \cdot X(\omega) < \sum_{\omega \notin B} P(\omega) \cdot 0.4m$.

So now we have $\mathbb{E}[X] < \sum_{\omega \in B} P(\omega) \cdot m + \sum_{\omega \notin B} P(\omega) \cdot 0.4m = mP(B) + 0.4mP(B^c)$, where B^c is the complementary event of B.

$P(B^c) = 1 - P(B)$, so $\mathbb{E}[X] < mP(B) + 0.4m(1 - P(B)) = 0.6mP(B) + 0.4m$.

But we assumed $P(B) < \frac{1}{6}$, thus $\mathbb{E}[X] < \frac{1}{6}0.6m + 0.4m = 0.5m$.

Finally we have $\mathbb{E}[X] < 0.5m$. But we have already shown $\mathbb{E}[X] = 0.5m$, thus we have reached a contradiction, and we can conclude $P(B) \geq \frac{1}{6}$. ∎

An improved algorithm By 4.1.1.3 we can formulate an improved algorithm which yields, by very good probability (specifically $1 - \frac{1}{e^{100}} \approx 99.999\%$), an assignment which satisfies at least 40% of the equations.

Notice that $\frac{1}{0.4} = 2.5$, so this is a high-probability 2.5-approximating algorithm. As opposed to the previous algorithm, here we don't say anything about the expectancy of the solutions the algorithms yields, but we say a much stronger statement about a specific solution. Altough 2.5-approximation is not

---

*Consider an equation of the form $a_1 x_1 + a_2 x_2 + ... + a_n x_n = b$. In F_2 each $x_i \in \{0, 1\}$. We can set $n - 1$ variables with either 1 or 0, by which the n'th variable is determined. This gives 2^{n-1} combinations.

as good as 2-approximation (40% vs. 50%), here we have, with high-probability, a specific solution which satisfies at least 40% of the equations.

Algorithm 4.2 Max-Lin-2 2.5-approximation high-probability algorithm

1. Choose 600 random assignments $\omega_1, \omega_2, ..., \omega_{600}$.

2. Return the assignment which satisfies the most equations.

By 4.1.1.3 we know that assignment ω_i satisfies 40% of the equations with probability of $\frac{1}{6}$. The probability that all 600 chosen assignments will fail to do this is $\left(1 - \frac{1}{6}\right)^{600}$. By lemma 5, we have $\left(1 - \frac{1}{6}\right)^{600} < e^{-100}$.

So the probability that at least one assignment will satisfy 40% of the equations is at least $1 - \frac{1}{e^{100}}$. ∎

4.1.2 3-SAT (Satisfiability)

See the definition of the problem in 3.1.7.

4.1.2.1 Probabilistic $\frac{7}{8}$-approximating algorithm

We show a probabilistic algorithm which gives a $\frac{7}{8}$-approximation with probability larger than $1 - \frac{1}{e^k}$ (with k as a parameter of the algorithm).

Algorithm 4.3 3-SAT probabilistic $\frac{7}{8}$-approximating algorithm

1. Randomly assign $\mathbb{T}$ or $\mathbb{F}$ to each variable x_i (with probability 0.5).

2. If $q(X) \geq \frac{7}{8}m$, finish and return X.

3. Repeat steps 1-2 at most $k(m+1)$ times.

4.1.2.2 Proof

We define *success* of the algorithm as the success in yielding a solution, X, which satisfies at least $\frac{7}{8}m$ clauses. In other words, we succeed only if $q(X) \geq \frac{7}{8}m$, otherwise we *fail*.

We will prove in two parts. First we will show that the probability for success of steps 1-2 (which will be reffered to as the *basic* steps) is $P(success\,basic) \geq \frac{1}{m+1}$. Then we will show that the probability for success of the entire algorithm (all 3 steps) is $P(success) \geq 1 - \frac{1}{e^k}$.

Part 1: We need to prove $P(success\,basic) \geq \frac{1}{m+1}$.

We define the probability space as all the tuples of $\mathbb{T}$ and $\mathbb{F}$ of size n: $\Omega = \{\mathbb{T}, \mathbb{F}\}^n$.

Obviously there are 2^n such tuples, so for some $\omega \in \Omega$ we have $P(\omega) = \frac{1}{2^n}$. We define two random variables:

- $X(\omega)$ - The amount of satisfied clauses by ω. Thus $P(success\,basic) = P\left(X \geq \frac{7}{8}m\right)$.

- $Y(\omega)$ - The amount of unsatisfied clauses by ω. Thus $P(failure\,basic) = P\left(Y > \frac{1}{8}m\right)$.

Instead of proving $P\left(X \geq \frac{7}{8}m\right) \geq \frac{1}{m+1}$, we can equivalently prove $P\left(Y > \frac{1}{8}m\right) < 1 - \frac{1}{m+1} = \frac{m}{m+1} = \frac{1}{1+\frac{1}{m}}$.

We define m new random variables:

$$\forall i = 1...m : Y_i(\omega) = \begin{cases} 1 & clause\,i\,is\,satisfied\,by\,\omega \\ 0 & clause\,i\,is\,unsatisfied\,by\,\omega \end{cases}.$$

Obviously $Y(\omega) = \sum_{i=1}^{m} Y_i(\omega)$, so by the linearity of the expectation $\mathbb{E}[Y] = \sum_{i=1}^{m} \mathbb{E}[Y_i]$.

$\mathbb{E}[Y_i] = P(Y_i = 1) \cdot 1 + P(Y_i = 1) \cdot 0 = P(Y_i = 1)$

$Y_i = 1$ means that clause i is satisfied. In 3-SAT, each clause contains 3 literals, each of them can evaluate to $\mathbb{T}$ or $\mathbb{F}$ by the assignment. Thus the probability for a clause to be satisfied is $\frac{1}{2} \cdot \frac{1}{2} \cdot \frac{1}{2} = \frac{1}{8}$. Thus $\mathbb{E}[Y_i] = P(Y_i = 1) = \frac{1}{8}$.

And now $\mathbb{E}[Y] = \sum_{i=1}^{m} \mathbb{E}[Y_i] = \frac{m}{8}$.

We now recall Markov's inequality:

For a non-negative random variable, Z, and some constant $c > 1$: $P(Z \geq c \cdot \mathbb{E}[Z]) \leq \frac{1}{c}$.

We use it on Y with $c = 1 + \frac{1}{m}$: $P\left(Y \geq \left(1 + \frac{1}{m}\right) \cdot \frac{m}{8}\right) \leq \frac{1}{1+\frac{1}{m}}$.

Both Y and m are integers, so $P\left(Y \geq \left(1 + \frac{1}{m}\right) \cdot \frac{m}{8}\right) = P\left(Y \geq \frac{m}{8} + \frac{1}{8}\right) = P\left(Y > \frac{m}{8}\right)$.

And finally we have $P(failure\,basic) = P\left(Y > \frac{m}{8}\right) \leq \frac{1}{1+\frac{1}{m}}$, as needed for this part of the proof. ■

Part 2: We need to prove $P(success) \geq 1 - \frac{1}{e^k}$.

Equivalently we can also prove $P(failure) < \frac{1}{e^k}$.

Failure in the entire algorithm is failure in all $k(m+1)$ attempts.

$P(failure) = \prod_{i=1}^{k(m+1)} P(failure\,in\,the\,i'th\,attempt)$.

In part 1 we already proved the failure in one attempt is $P(failure\,basic) \leq \frac{1}{1+\frac{1}{m}} = 1 - \frac{1}{m+1}$, thus:

$$P(failure) \leq \prod_{i=1}^{k(m+1)} \left(1 - \frac{1}{m+1}\right) = \left(1 - \frac{1}{m+1}\right)^{k(m+1)} = \left(\left(1 - \frac{1}{m+1}\right)^{(m+1)}\right)^k <$$

$\left(\frac{1}{e}\right)^k = \frac{1}{e^k}$, thus completing the proof. ■

Chapter 5

Flow Networks

5.1 Definitions and properties

Definition 19. Flow network.
A flow network is a tuple of 5 elements (V, E, C, s, t):

- $G = (V, E)$ - a **directed** graph with vertices V and edges E.

- $C : E \to \mathbb{R}^+$ - a capacity function on the edges of the network.

- $s \in V$ - the source vertex.

- $t \in V$ - the terminal vertex.

Definition 20. Flow.
A flow in a flow network is a function $f : E \to \mathbb{R}^+$ which maintains:

1. The flow in an edge cannot exceed the capacity of the edge: $\forall e \in E : f(e) \leq C(e)$.

2. Mass preservation for all inner vertices:
$$\forall v \in V \setminus \{s, t\} : \underbrace{\sum_{u \in V : (u,v) \in E} f(u, v)}_{incoming\ flow\ into\ v} = \underbrace{\sum_{u \in V : (v,u) \in E} f(v, u)}_{outgoing\ flow\ from\ v}.$$

Definition 21. Flux.
The flux of a flow is the net flow coming out of the source s:
$$|f| = \underbrace{\sum_{u \in V : (s,u) \in E} f(s, u)}_{outgoing\ flow\ from\ s} - \underbrace{\sum_{u \in V : (u,s) \in E} f(u, s)}_{incoming\ flow\ into\ s}.$$

Definition 22. Rational flow.
f is a rational flow if there are no two vertices $u, v \in V$ such that both $f(u, v) > 0$ and $f(v, u) > 0$.

Definition 23. Optimal flow.

Usually when working with flow networks, the goal is to find the maximal flux legal flow in the network. If f has the maximal flux in the network, it is said to be optimal.

Lemma 24. *For every network there is an optimal flow which is also rational.*

Proof. Let f be an optimal flow in the network with a minimal number of pairs $u, v \in V$ such that both $f(u, v) > 0$ and $f(v, u) > 0$. We will show that the number of these pairs is 0, and so f is rational.

We prove by negation.

Let's assume that there is a pair $u, v \in V$ such that both $f(u, v) > 0$ and $f(v, u) > 0$. We will build a new flow, g, and will show that it is also optimal, and has less "bad" pairs than in f, thus reaching a contradiction (because we assumed f is the optimal flow with minimal number of "bad" pairs).

We assume, WLOG, that $f(u, v) > f(v, u)$.

$$g(e = (w, z)) = \begin{cases} f(e) & \{w, z\} \neq \{u, v\} \\ f(u, v) - f(v, u) & w = u, z = v \\ f(v, u) - f(v, u) = 0 & w = v, z = u \end{cases}$$

We now need to show that g is legal, optimal and has less "bad" pairs than f.

First we notice $\forall e \in E : g(e) \leq f(e) \leq C(e)$.

Now we check mass preservation on g. The only vertices that were changed are u, v. For both u, v we decreased the incoming and outgoing flow by the same amount ($f(v, u)$), thus mass is still preserved. So g is legal.

Now we check g is optimal.

$|g| =$

TODO $\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\square$

From now on, when we talk about flows, we will mean rational flows.

We would like to work with flow networks that are on complete graphs. So for every given flow network, we can define an *extended* flow network, which is completely equivalent.

Definition 25. Expanded capacity.

We define a capacity on the complete graph (for **every** two vertices in V):

$$C'(u, v) = \begin{cases} C(u, v) & (u, v) \in E \\ 0 & (u, v) \notin E \end{cases}.$$

Definition 26. Expanded flow.

We define a flow on the complete graph (for **every** two vertices in V):

$$f'(u, v) = \begin{cases} f(u, v) & f(u, v) > 0 \\ -f(u, v) & f(v, u) > 0 \\ 0 & otherwise \end{cases}.$$

Notice that a normal flow (as defined before) is always non-negative, but an exapnded flow can be negative $\left(f' : V \times V \to \mathbb{R}\right)$.

We also notice that a legal expanded flow holds the following properties:

1. Anti-symmetricity: $\forall u, v \in V : f'(u, v) = -f'(v, u)$.

2. Constrained by expanded capacity: $\forall u, v \in V : f'(u, v) \leq C'(u, v)$.

3. Preserves mass on inner vertices: $\forall u \in V \setminus \{s, t\} : \sum_{v \in V} f'(u, v) = 0$.

Lemma 27. *If f is a legal flow, then the expanded f' is a legal expanded flow.*

Proof. We will prove for all three properties:
TODO $\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad\qquad$ $\square$

5.2 Ford and Fulkerson's algorithm

Here we present Ford & Fulkerson's algorithm for finding the optimal flow in a flow network.

First we start with some definitions.

Definition 28. Residual flow network.
Let $G = (V, E, C, s, t)$ be a flow network, and f a legal flow in it.
We define G's and f's residual network $G_f = (V, E_f, C_f, s, t)$ where:

- $C_f(u, v) = C(u, v) - f(u, v)$ (Residual capacity)

- $E_f = \{(u, v) : C_f(u, v) > 0\}$

Definition 29. Augmentation path.
An augmentation path of flow f is a simple path (which visits each vertex at most once), in the resudual network G_f, from s to t.

Definition 30. Residual capacity of an augmentation path.
Let P be an augmentation path in G_f. We define the residual capacity of P as the minimal capacity of all the edges of P: $\Delta_P = \min_{e \in P} \{C_f(e)\}$

Definition 31. Residual flow along the augmentation path.
This is the maximal flow we can enforce on the augmentation path without violating any constraints (capacities, mass preservation, etc.).
We define $\Delta_P f(u, v) = \begin{cases} \Delta_P & (u, v) \in P \\ -\Delta_P & (v, u) \in P \\ 0 & otherwise \end{cases}$.
We notice that $\Delta_P f$ is a legal flow in G_f.

Algorithm 5.1 Ford and Fulkerson's algorithm

1. Initialize $f = 0$.

2. While there is an augmentation path in G_f:

   (a) Find some augmentation path in G_f and denote it with P.

   (b) Update: $f = f + \Delta_P f$.

3. Return f.

5.2.1 Proof

We need to show:

1. $f + \Delta_P f$ is a legal flow.

2. In each iteration the flux is improved: $|f + \Delta_P f| > |f|$.

3. As long as there is an augmentation path in the residual network of the current flow, then it is not optimal.

5.2.1.1 Legality

We will show that $f + \Delta_P f$ is a legal flow in the flow network.
   We denote $g = f + \Delta_P f$.

Anti-symmetricity   Both f and $\Delta_P f$ are anti-symmetric, and so is g.

Capacity constraints $g(u,v) = f(u,v) + \Delta_P f(u,v) \le f(u,v) + C_f(u,v) = f(u,v) + C(u,v) - f(u,v) = C(u,v)$, as needed.

Mass preservation $\forall u \in V \setminus \{s,t\} : \sum_{v \in V} g(u,v) = \sum_{v \in V} (f(u,v) + \Delta_P f(u,v)) = \underbrace{\sum_{v \in V} f(u,v)}_{=0} + \underbrace{\sum_{v \in V} \Delta_P f(u,v)}_{=0} = 0$, as needed.

∎

5.2.1.2 Flux improvement

$|f + \Delta_P f| = |f| + |\Delta_P f| > |f|$

∎

5.2.1.3 Optimallity (Max-Flow Min-Cut theorem)

We want to show that as long as there is an augmentation path in the residual network of the current flow, then it is not optimal.

We will prove a stronger statement:

Theorem 32. *Max-Flow Min-Cut.*
The following are equivalent:

1. f is an optimal flow in network G.

2. G_f has no augmentations paths.

3. There exists a cut (S, T) such that $f(S, T) = C(S, T)$.

Proof. We will prove $1 \to 2$, $2 \to 3$, $3 \to 1$

$1 \to 2$:

f is an optimal flow in network G and we want to show that G_f has no augmentations paths.

If there was an augmentation paths we could have augmented the flow by $f = f + \Delta_P f$ and improve the flow. But this contradicts the optimality of f. Thus indeed there is no augmentation path in G_f.

$2 \to 3$:

G_f has no augmentations paths and we want to show that there exists a cut (S, T) such that $f(S, T) = C(S, T)$.

G_f has no augmentations paths, which means there is no path from s to t in the graph (V, E_f). We can define the following cut (S, T):

S - every vertex $v \in V$ such that there is a path from s to v.

$T = V \backslash S$.

Obviously $s \in S$ and $t \in T$. We notice that there are no edges of E_f in the cut. I.e. $\forall u \in S, v \in T : (u, v) \notin E_f$, otherwise there was a path from s to $v \in T$, which contradicts the definition of S and T. By the definition of E_f this means that $\forall u \in S, v \in T : C_f(u, v) = 0$. And by the definition of C_f we have $\forall u \in S, v \in T : C(u, v) = f(u, v)$. So now we can conclude $f(S, T) = \sum_{u \in S, v \in T} f(u, v) = \sum_{u \in S, v \in T} C(u, v) = C(S, T)$, as needed.

$3 \to 1$:

There exists a cut (S, T) such that $f(S, T) = C(S, T)$ and we want to show that f is an optimal flow in network G.

We have already seen in [TODO] that for any cut $\left(S', T'\right)$ and any flow g in the network it holds that $|g| \leq C\left(S', T'\right)$. If f was not optimal, then there is some other flow g such that $|g| > |f|$. But then, for any cut $\left(S', T'\right)$ we would have $|f| < C\left(S', T'\right)$, and specifically for (S, T): $|f| < C(S, T)$. This is in contradiction to $f(S, T) = C(S, T)$. Thus f is indeed optimal. $\qquad \square$

5.3 Edmonds and Karp's algorithm

Edmond & Karp's algorithm is actually an improvement to Ford & Fulkerson's algorithm. In step 2a of algorithm 5.1, instead of choosing some arbitrary augmentation path, we choose the shortest augmentation path (length wise).

Algorithm 5.2 Edmonds and Karp's algorithm

1. Initialize $f = 0$.

2. While there is an augmentation path in G_f:

   (a) Find the shortest augmentation path in G_f and denote it with P.

   (b) Update: $f = f + \Delta_P f$.

3. Return f.

In order to find the shortest augmentation path, we can use, for example, the BFS algorithm. The runtime of BFS is $O(|E|)$.

5.3.1 Edmond and Karp's theorem

Theorem 33. *Edmonds and Karp.*
Algorithm 5.2 performs $O(|E| \cdot |V|)$ iterations.

5.3.1.1 Proof

First, some denotations:

- f_i the flow after the i'th iteration.

- P_{i-1} the augmentation path of step 2a in the i'th iteration.

- $G_{f_i} = (V, E_{f_i}, C_{f_i}, s, t)$ the residual network after the i'th iteration.

- $\delta_i(v)$ the minimal distance of vertex $v \in V$ from the source s in the graph (V, E_{f_i}).

Some remarks:

Remark 34. $f_i = f_{i-1} + \Delta_{P_{i-1}} f_{i-1}$.

Remark 35. The algorithm stops when $\delta_i(t) = \infty$, because this means there is no augmentation path in G_{f_i} (no way to reach t from s).

Remark 36. Let P be a path of minimal length between s and some vertex v. Then for every vertex x in this path, the distance from s to x in this path is the minimal distance between s and x in the graph. Specifically, if x and y are two adjacent vertices in the path such that x appears before y, then $\delta(y) = \delta(x) + 1$.

Remark 37. In the i'th iteration, the flow f_i is larger than f_{i-1} in all the edges along the augmentation path P_{i-1} by $\Delta_{P_{i-1}}$, and it is smaller than f_{i-1}, in the opposite direction, by $\Delta_{P_{i-1}}$ (we conclude this from definition 31).

Corollary 38. *If $(u,v) \in G_{f_{i-1}}$ and $(u,v) \notin G_{f_i}$ then the flow in (u,v) was reduced, which must mean that the flow in (v,u) was increased, thus $(v,u) \in P_{i-1}$.*

Some lemmas:

Lemma 39. *$\forall v \in V : \delta_i(v) \leq \delta_{i+1}(v) \leq \infty$ and if $\delta_i(v) = \infty$ then also $\delta_{i+1}(v) = \infty$.*

Proof. First we will show the second part.

We prove by negation. Assume there is a vertex $v \in V$ and an index i for which $\delta_i(v) = \infty$ but $\delta_{i+1}(v) < \infty$. If there are several such vertices, we choose v to be the closest one to s in $G_{f_{i+1}}$. We denote the shortest path from s to v in $G_{f_{i+1}}$ with P, and the vertex which comes before v in P as u, so $P = (s, ..., u, v)$. Because $\delta_{i+1}(v) < \infty$ then also $\delta_{i+1}(u) < \infty$ (because u is closer to s than v). And because $\delta_{i+1}(u) < \infty$ then also $\delta_i(u) < \infty$ (we chose v as the closest vertex to s which doesn't hold this). Now we have $\delta_i(v) = \infty$ and $\delta_i(u) < \infty$, which means that $(u,v) \notin G_{f_i}$. But $(u,v) \in G_{f_{i+1}}$, so by corollary 38 the edge (v,u) was in the augmentation path of the i'th iteration, which means that there is a path from s to v in the i'th iteration. This contradicts the assumption that $\delta_i(v) = \infty$. Thus we conclude that if $\delta_i(v) = \infty$ then also $\delta_{i+1}(v) = \infty$.

Now we will show the first part.

Again we prove by negation. Assume there is a vertex $v \in V$ and an index i for which $\delta_i(v) > \delta_{i+1}(v)$. If there are several such vertices, we choose v to be the closest one to s in $G_{f_{i+1}}$. We denote the shortest path from s to v in $G_{f_{i+1}}$ with P, and the vertex which comes before v in P as u, so $P = (s, ..., u, v)$. It is evident that $\delta_{i+1}(v) = \delta_{i+1}(u) + 1$. Also, $\delta_i(u) \leq \delta_{i+1}(u)$ (we chose v as the closest vertex to s which doesn't hold this). Now we have $(u,v) \notin G_{f_i}$, because otherwise $\delta_{i+1}(v) = \delta_{i+1}(u) + 1 \geq \delta_i(u) + 1 \underbrace{=}_{if\ (u,v)\in G_{f_i}} \delta_i(v)$, which

contradicts our assumption. So finally we have $(u,v) \notin G_{f_i}$ and $(u,v) \in G_{f_{i+1}}$. As before, we conclude that (v,u) was an edge in the augmentation path in the i'th iteration, so $\delta_i(u) = \delta_i(v) + 1$. Remember the assumption $\delta_i(v) > \delta_{i+1}(v)$. So now $\delta_i(v) > \delta_{i+1}(v) = \delta_{i+1}(u) + 1 \geq \delta_i(u) + 1 = \delta_i(v) + 2$. Finally we have $\delta_i(v) > \delta_i(v) + 2$, which is obviously impossible. Thus we conclude that $\forall v \in V : \delta_i(v) \leq \delta_{i+1}(v)$. $\square$

Definition 40. Critical edge.

(u,v) is a critical edge for flow f_{i-1} if its residual capacity $C_{f_{i-1}}(u,v)$ in the augmentation path P_{i-1} is minimal. In other words, this edge is the bottleneck of P_{i-1}.

Lemma 41. *For every $u,v \in V$ the edge (u,v) can be a critical edge in at most $O(|V|)$ iterations.*

Proof. Let i, j $(i < j)$ be two different iterations in algorithm 5.2 such that in both the edge (u, v) is a critical edge.

Because (u, v) is critical in the i'th iteration, it means that $(u, v) \notin G_{f_i}$. But by the assumption, there is some later iteration in which (u, v) is again a critical edge. This means that (u, v) must reappear in some iteration k between i and j $(i < k < j)$. This means that at some point there was a flow in the opposite direction (on (v, u)), in other words $(v, u) \in P_k$. This means that in the k'th iteration $\delta_k(u) = \delta_k(v) + 1$. Also, because (u, v) is an edge in the augmentation path of the i'th iteration we have $\delta_i(v) = \delta_i(u) + 1$.

From lemma 39 we have $\delta_k(v) \geq \delta_i(v)$ and $\delta_j(u) \geq \delta_k(u)$. Thus $\delta_j(u) \geq \delta_k(u) = \delta_k(v) + 1 \geq \delta_i(v) + 1 = \delta_i(u) + 1 + 1 = \delta_i(u) + 2$. Finally we obtained $\delta_j(u) \geq \delta_i(u) + 2$.

Obviously $\delta(u)$ is bounded by $|V|$ (for any iteration, as long as $\delta(u) < \infty$). Because $\delta(u)$ increases by at least 2 between any two iterations in which (u, v) is critical, we conclude that (u, v) can be critical at most $\frac{|V|}{2}$ times, which is $O(|V|)$, as needed. $\qquad\square$

Corollary 42. *Because there are at most $2|E|$ edges in each residual network, and because each edge can be critical $O(|V|)$ times (by lemma 41), we conclude that the algorithm performs $O(|E| \cdot |V|)$ iterations.*

With corollary 42 we complete the proof. $\blacksquare$

5.4 Examples

5.4.1 Maximal matching in a bipartite graph

We have an undirected bipartite graph $G = (L, R, E)$, and we want to find a maximal matching.

5.4.1.1 Algorithm which uses a flow network

Algorithm 5.3 Maximal matching in a bipartite graph (with flow network)

---

1. Construct a flow network $G = \left(V^{'}, E^{'}, C, s, t\right)$:

   (a) $V^{'} = L \cup R \cup \{s, t\}$ (with s and t as some artificial vertices).

   (b) $E^{'} = \{(s, v) : v \in L\} \cup \{(v, t) : v \in R\} \cup \vec{E}$ (where $\vec{E}$ is a set of the edges in E directed from L to R).

   (c) $\forall e \in E^{'} : C(e) = 1$.

2. Find the maximal flow on the network.

3. Return all the original edges on which the flow is 1: $M = \left\{e \in \vec{E} | f(e) = 1\right\}$.

5.4.1.2 Proof of legality

We need to show that the solution of the algorithm (M) is indeed a legal matching.

   We prove by negation. Assume there is a vertex v which touces more than one edge in M.

   If $v \in R$, then it has one outgoing edge (v, t) which has capacity $C(v, t) = 1$. Because all the capacities are integers...

   [TODO]

5.4.2 Roads and advertisements

A town has a system of 1-way or 2-way roads and junctions. We denote $A = (a_1, a_2, ..., a_n)$ the set of junctions. $E \subseteq A$ is a subset of entrance junctions, through which one can enter the town. We have a store on one of the junctions which we want to advertise. We want our advertisements to appear on at least one junction on any path from one of the entrance junctions to our store's junction. The cost of putting an advertisement on junction a_i is s_i. Our goal is to find a subset of the junctions $A^{'} \subseteq A$ that satisfies this condition, and that has a minimal cost (minimize $\sum_{a_i \in A'} s_i$).

5.4.2.1 Algorithm which uses a flow network

TODO

Chapter 6

Fast Fourier Transform

6.1 Proof that FFT solves DFT

6.1.1 Denotations and definitions

- Coefficients representation of a polynomial: The vector which represents the polynomial in the monomial basis $1, x, x^2, ..., x^{n-1}$ is called the *coefficients representation*. If we have coefficients representation $(a_0, a_1, ..., a_{n-1})$ then the polynomial is given by $p(x) = \sum_{i=0}^{n-1} a_i x^i$.

- Values representation of a polynomial: Given n points $x_0, x_1, ..., x_{n-1}$, the values of the polynomial in these points $((p(x_0), p(x_1), ..., p(x_{n-1})))$ is called the *values representation* of the polynomial. The vector of these values is in the Lagrange basis $L_0(x), L_1(x), ..., L_{n-1}(x)$, where $L_i(x) = \frac{(x-x_0)\cdot...\cdot(x-x_{i-1})\cdot(x-x_{i+1})\cdot...\cdot(x-x_{n-1})}{(x_i-x_0)\cdot...\cdot(x_i-x_{i-1})\cdot(x_i-x_{i+1})\cdot...\cdot(x_i-x_{n-1})}$. The polynomial is given by $p(x) = \sum_{i=0}^{n-1} p(x_i) L_i(x)$.

- ω_n - primitive root of unity of order n: $\omega_n = e^{\frac{2\pi i}{n}}$.

- M - the Vandermonde matrix which we use for the basis transformation from the monomial basis to the Lagrange basis. If the values representation of the polynomial is needed on points $x_0, x_1, ..., x_{n-1}$, then $M = \begin{pmatrix} x_0^0 & x_0^1 & \cdots & x_0^{n-1} \\ x_1^0 & x_1^1 & \cdots & x_1^{n-1} \\ \vdots & \vdots & \ddots & \vdots \\ x_{n-1}^0 & x_{n-1}^1 & \cdots & x_{n-1}^{n-1} \end{pmatrix} = \left(x_i^j\right)_{i,j=0}^{n-1}$. And now $\begin{pmatrix} p(x_0) \\ p(x_1) \\ \vdots \\ p(x_{n-1}) \end{pmatrix} = M \begin{pmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{pmatrix}$. If we choose the points to be the roots of unity of order n,

$1, \omega_n, \omega_n^2, ..., \omega_n^{n-1}$, then we have $M = \left((\omega_n)^{i \cdot j} \right)_{i,j=0}^{n-1}$.

- M^{-1} - the inverse Vandermonde matrix which we use for the basis transformation from the Lagrange basis to the monomial basis. If we choose the points to be the roots of unity of order n, $1, \omega_n, \omega_n^2, ..., \omega_n^{n-1}$, then we have $M^{-1} = \frac{1}{n} \left((\omega_n)^{-i \cdot j} \right)_{i,j=0}^{n-1}$.

6.1.2 The lemma

Given a coefficients representation $\begin{pmatrix} a_0 \\ \vdots \\ a_{n-1} \end{pmatrix}$ we want to calculate the corresponding values representation **on the roots of unity** $\begin{pmatrix} p\left((w_n)^0\right) \\ p\left((w_n)^1\right) \\ \vdots \\ p\left((w_n)^{n-1}\right) \end{pmatrix}$.

We already know that, given the corresponding Vandermonde matrix, $\begin{pmatrix} p\left((w_n)^0\right) \\ p\left((w_n)^1\right) \\ \vdots \\ p\left((w_n)^{n-1}\right) \end{pmatrix} =$

$M \begin{pmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{pmatrix}$. But this is a lengthy calculation - multiplying a matrix with a vector is $O\left(n^2\right)$. We want to do this in $O\left(n \log n\right)$.

We denote the matrix M with DFT_n and M^{-1} with DFT_n^{-1}.

Lemma 43. *We can make the calculation shorter with the following recurrsion formula:*

$$DFT_n \begin{pmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{pmatrix} = \begin{pmatrix} DFT_{\frac{n}{2}} \begin{pmatrix} a_0 \\ a_2 \\ \vdots \\ a_{n-2} \end{pmatrix} \\ DFT_{\frac{n}{2}} \begin{pmatrix} a_0 \\ a_2 \\ \vdots \\ a_{n-2} \end{pmatrix} \end{pmatrix} + \begin{pmatrix} (w_n)^0 \\ (w_n)^1 \\ \vdots \\ (w_n)^{n-1} \end{pmatrix} \cdot \begin{pmatrix} DFT_{\frac{n}{2}} \begin{pmatrix} a_1 \\ a_3 \\ \vdots \\ a_{n-1} \end{pmatrix} \\ DFT_{\frac{n}{2}} \begin{pmatrix} a_1 \\ a_3 \\ \vdots \\ a_{n-1} \end{pmatrix} \end{pmatrix}$$

(with the last product $\cdot$ as an element-by-element product).

Remark 44. Lemma 43 works for any n (the order of the roots of unity) which is a power of 2. But this can always be arranged by padding the coefficients representation of the polynomial with zeros in the end.

6.1.3 The proof

We will now prove lemma 43. We do this by showing that the vectors on both sides of the equation are equal on all indices.

We denote $m = \frac{n}{2}$ (remember n is a power of 2), ω as the primitive root of unity of order n and η as the primitive root of unity of order m. Notice that $\eta = \omega^2$.

We also denote $M = DFT_n = \left(\omega^{a \cdot b}\right)_{a,b=0}^{n-1}$ and $L = DFT_m = \left(\eta^{a \cdot b}\right)_{a,b=0}^{m-1}$.

We will now examine the k'th element on the vectors in both sides of the equation.

Case 1 - $0 \leq k \leq m-1$ $\left(0 \leq k \leq \frac{n}{2}-1\right)$: $DFT_n \begin{pmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{pmatrix}$ in the k'th

element is given by $\sum_{t=0}^{n-1} [M]_{k,t} \cdot a_t = \sum_{t=0}^{n-1} \omega^{k \cdot t} \cdot a_t$.

We can decompose this expression as follows:

$\sum_{t=0}^{n-1} \omega^{k \cdot t} \cdot a_t = \underbrace{\sum \omega^{k \cdot t} \cdot a_t}_{even\ t's} + \underbrace{\sum \omega^{k \cdot t} \cdot a_t}_{odd\ t's}$

- $\underbrace{\sum}_{even\ t's} \omega^{k \cdot t} \cdot a_t = a_0 + \omega^{2k} a_2 + \omega^{4k} a_4 + ... + \omega^{(n-2)k} a_{n-2} = a_0 + \eta^k a_2 +$

$\eta^{2k} a_4 + ... + \eta^{(m-1)k} a_{n-2} = \sum_{t=0}^{m-1} \eta^{k \cdot t} \cdot a_{2t} = \sum_{t=0}^{m-1} [L]_{k,t} \cdot a_{2t}$, which is

precisely the k'th element in $\begin{pmatrix} DFT_{\frac{n}{2}}\begin{pmatrix} a_0 \\ a_2 \\ \vdots \\ a_{n-2} \end{pmatrix} \\ DFT_{\frac{n}{2}}\begin{pmatrix} a_0 \\ a_2 \\ \vdots \\ a_{n-2} \end{pmatrix} \end{pmatrix}$.

- $\underbrace{\sum}_{odd\ t's} \omega^{k \cdot t} \cdot a_t = \omega^k a_1 + \omega^{3k} a_3 + \omega^{5k} a_5 + ... + \omega^{(n-1)k} a_{n-1} = \omega^k \left(a_1 + \omega^{2k} a_3 + \omega^{4k} a_5 + ... + \omega^{(n-2)k} a_{n-1}\right) =$

$\omega^k \left(a_1 + \eta^k a_3 + \eta^{2k} a_5 + ... + \eta^{(m-1)k} a_{n-1}\right) = \omega^k \left(\sum_{t=0}^{m-1} \eta^{k \cdot t} \cdot a_{2t+1}\right) =$

$\omega^k \sum_{t=0}^{m-1} [L]_{k,t} \cdot a_{2t+1}$ which is precisely the k'th element in $\begin{pmatrix} (w_n)^0 \\ (w_n)^1 \\ \vdots \\ (w_n)^{n-1} \end{pmatrix}$.

$$\begin{pmatrix} DFT_{\frac{n}{2}} \begin{pmatrix} a_1 \\ a_3 \\ \vdots \\ a_{n-1} \end{pmatrix} \\ DFT_{\frac{n}{2}} \begin{pmatrix} a_1 \\ a_3 \\ \vdots \\ a_{n-1} \end{pmatrix} \end{pmatrix}.$$

Finally we have that the k'th element in the left side of the equation equals the k'th element in the right side of the equation.

Case 2 - $m \leq k \leq 2m-1$ $\left(\frac{n}{2} \leq k \leq n-1\right)$: Again, $DFT_n \begin{pmatrix} a_0 \\ a_1 \\ \vdots \\ a_{n-1} \end{pmatrix}$ in the k'th element is given by $\sum_{t=0}^{n-1} [M]_{k,t} \cdot a_t = \sum_{t=0}^{n-1} \omega^{k \cdot t} \cdot a_t$.

Now we can write $k = m + s$ and $0 \leq s \leq m - 1$. We notice that $\eta^k = \eta^{m+s} = \eta^m \cdot \eta^s = 1 \cdot \eta^s = \eta^s$.

We can now repeat a similar analysis as in case 1 and reach the conclusion that the k'th element in the left side of the equation equals the $(m + s)$'th (or, k'th) element in the right side of the equation.

Thus we complete the proof. ∎

Chapter 7

Cryptography

7.1 Foundations

7.1.1 Survey of group theory

Definition 45. Group.

A non-empty group G with some operation $\circ$, such that $\circ : G \times G \to G$, holds the following properties:

1. Associativity: $\forall a, b, c \in G : a \circ (b \circ c) = (a \circ b) \circ c$.

2. Existence of *identity element* e: There exists $e \in G$ such that for every $a \in G$ it holds $a \circ e = e \circ a = a$.

3. Existence of *inverse element:* For every $a \in G$ there exists $a^{-1} \in G$ such that $a \circ a^{-1} = a^{-1} \circ a = e$.

Remark 46. We denote a group G with operation $\circ$ as a tuple $(G, \circ)$.

Lemma 47. *The identity element of G is unique.*

Proof. Assume there are two identity elements a and b.
Because a is an identity element we have $a \circ b = b$.
Because b is an identity element we have $a \circ b = a$.
Thus we conclude $a = b$, which means the identity element is unique. $\qquad\square$

Lemma 48. *For each $a \in G$ there is a unique inverse element $a^{-1} \in G$.*

Proof. Assume a has two inverse elements b, c.
Because c is an inverse element of a: $b \circ (a \circ c) = b \circ e = b$.
Because b is an inverse element of a: $b \circ (a \circ c) = (b \circ a) \circ c = e \circ c = c$.
Thus we conclude $b = c$, which means the inverse element is unique. $\qquad\square$

[TODO - more definitions. such as sub groups, $ord_G(a)$, etc.]

Theorem 49. *Lagrange's theorem.*

Let G be a finite group and $H \subseteq G$ a subgroup, then $|H|$ divides $|G|$ ($|G|$ mod $|H| = 0$).

Proof. For each $a \in G$ we denote $aH = \{a \cdot h | h \in H\}$.

We will see that for every $a, b \in G$ either $aH = bH$ or $aH \cap bH = \emptyset$.

Let's assume there is $x \in aH \cap bH$. It means that there are $h_1, h_2 \in H$ such that $x = a \cdot h_1 = b \cdot h_2$. Thus $b^{-1} \cdot a = h_2 \cdot h_1^{-1}$ which means that $b^{-1} \cdot a \in H$.

We will now show that $aH \subseteq bH$ and $bH \subseteq aH$, and conclude that $aH = bH$.

We will show $aH \subseteq bH$, and the other direction is analogous.

Let $x = bh \in bH$. We can write $x = b\underbrace{\left(\left(b^{-1}a \right) \left(b^{-1}a \right)^{-1} \right)}_{e} h = b \left(b^{-1}a \right) \left(b^{-1}a \right)^{-1} h = a \left(b^{-1}a \right)^{-1} h$.

We have seen that $b^{-1}a \in H$, and thus also $\left(b^{-1}a \right)^{-1} \in H$ and $\left(b^{-1}a \right)^{-1} h \in H$. Thus we conclude $x = a \left(b^{-1}a \right)^{-1} h \in aH$.

We have seen that if $x \in bH$ then $x \in aH$, which means $bH \subseteq aH$, as needed (along with the opposite direction of $aH \subseteq bH$) to conclude that $aH = bH$.

Now we will see that for every $a, b \in G$: $|aH| = |bH|$. Obviously it's true if $aH = bH$, so we assume that $aH \cap bH = \emptyset$. We will define a bijection $\varphi : aH \to bH$. If we show that φ is indeed a bijection, then we can conclude $|aH| = |bH|$.

We define $\varphi(x) = b \cdot a^{-1}x$.

For $x = a \cdot h \in aH$ we have $\varphi(x) = b \cdot a^{-1}a \cdot h = b \cdot h \in bH$.

First we show that φ is injective.

For some x, y we have $\varphi(x) = \varphi(y) \Leftrightarrow b \cdot a^{-1}x = b \cdot a^{-1}y \Leftrightarrow \left(b^{-1} \cdot a \right) \cdot b \cdot a^{-1}x = \left(b^{-1} \cdot a \right) \cdot b \cdot a^{-1}y \Leftrightarrow x = y$. So φ is indeed injection.

Now we will show that φ is surjective.

Each $y \in bH$ has the form $y = b \cdot h$ which is exactly $\varphi(x)$ with $x = a \cdot h \in aH$. Thus φ is surjective.

we have shown that φ is indeed a bijection and we can conclude $|aH| = |bH|$. [TODO] $\qquad\square$

7.1.2 Greatest common divisor (GCD)

Definition 50. Greatest Common Divisor.

Let $a, b \in \mathbb{N}$ such that $a \geq b$. The GCD of a and b is the largest $d \in \mathbb{N}$ ($d \geq 1$) such that d divides both a and b without a remainder.

$gcd(a, b) = \max \{d \in \mathbb{N}, d \geq 1 : a \mod d = b \mod d = 0\}$

Theorem 51. *Fundamental theorem of arithmetic.*

Every $n \in \mathbb{N}$ ($n > 1$) can be written as a product of prime numbers. The primes involved in the product are unique.

$n = p_1^{r_1} \cdot p_2^{r_2} \cdot \ldots \cdot p_k^{r_k} = \prod_{i=1}^{k} p_i^{r_i}$, *where each p_i is a prime number and each $e_i \in \mathbb{N} \cup \{0\}$.*

Definition 52. Greatest Common Divisor (alternative).

Let $a, b \in \mathbb{N}$ such that $a \geq b > 1$.

By theorem 51 we can write:

$a = p_1^{e_1} \cdot p_2^{e_2} \cdot ... \cdot p_k^{e_k}$

$b = p_1^{f_1} \cdot p_2^{f_2} \cdot ... \cdot p_k^{f_k}$

Where $\{p_i\}_{i=1}^k$ is the set of prime numbers which devide both a and b (remember e_i and f_i can be zero).

Then: $gcd(a, b) = p_1^{\min\{e_1, f_1\}} \cdot p_2^{\min\{e_2, f_2\}} \cdot ... \cdot p_k^{\min\{e_k, f_k\}} = \prod_{i=1}^k p_i^{\min\{e_i, f_i\}}$.

Claim 53. Let $a, b \in \mathbb{N}$ such that $a \geq b$, and $S = \{a \cdot x + b \cdot y | x, y \in \mathbb{Z}, a \cdot x + b \cdot y \geq 1\}$.

Then: $gcd(a, b) = \min S$.

Proof. We denote $t = gcd(a, b)$ and $z = \min S$. We want to prove $t = z$. For this we will show that $t \leq z$ and $z \leq t$ and conclude $t = z$.

$t \leq z$:

We will see that $z \mod t = 0$, which means that $t \leq z$.

$t = gcd(a, b)$ so $a \mod t = 0$ and $b \mod t = 0$, which means that there exist some $k_1, k_2 \in \mathbb{N}$ such that $t \cdot k_1 = a$ and $t \cdot k_2 = b$.

Also, $z \in S$, which means there exist some $x, y \in \mathbb{Z}$ such that $z = ax + by$.

Now we can write $z = ax + by = t \cdot k_1 \cdot x + t \cdot k_2 \cdot y = t(k_1 x + k_2 y)$. Which means that $z \mod t = 0$, as needed.

$z \leq t$:

We will show that z divides both a and b, but because t is the maximal such divisor, we will conclude $z \leq t$.

We will prove for a, and the proof for b is analogous.

First we will show that $z \leq a$. Assume it isn't, so $z > a$. We define $z' = 1 \cdot a + 0 \cdot b$. By definition of S, $z' \in S$, and obviously $z' = a$. Thus $z' < z$, but this contradicts the minimality of z in S. So $z \leq a$.

Now we can write $a = k \cdot z + r$ with $0 \leq r \leq z - 1$. All that is left is to show that $r = 0$ (and thus conclude that z divides a).

$z \in S$ so it must be of the form $z = ax + by$, which means $r = a - k \cdot z = a(1 - x) - by$. If $r \geq 1$, then also $r \in S$ (it fits the structure of S), and from the minimality of z in S we conclude $r \geq z$. But r was defined $r \leq z - 1$, so we have a contradiction, and thus $r = 0$, as needed. $\square$

7.1.2.1 Euclid's algorithm for GCD

Index

List of Figures